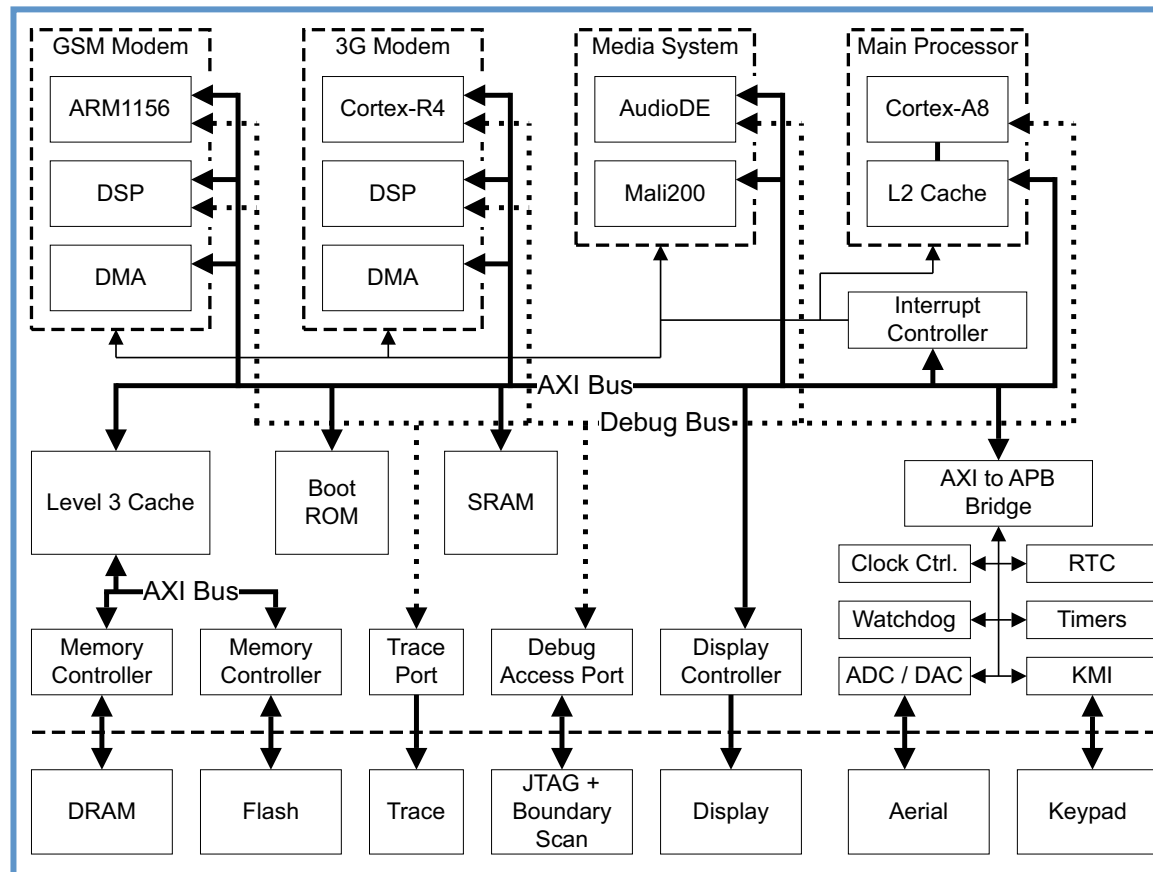


# Hardware and Booting

Chester Rebeiro  
IIT Madras




# System Organization




# Address Types

- Memory Addresses
- Memory Mapped IO Addresses

# Memory Maps



# Accessing keyboards



2 input ports and 1 output port

Pointer to the keyboard:

```
char *keyboard_ptr = (char *) 0x10000000;
```

Read from keyboard


```
char x = *(keyboard_ptr + 1);
```

Write to keyboard



```
*(keyboard) = y;
```

# Powering Up

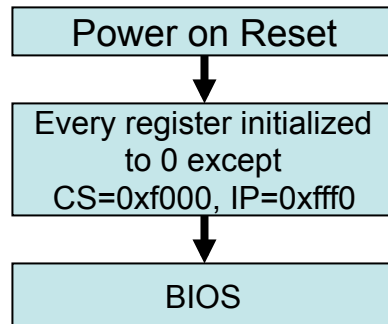
Power on Reset



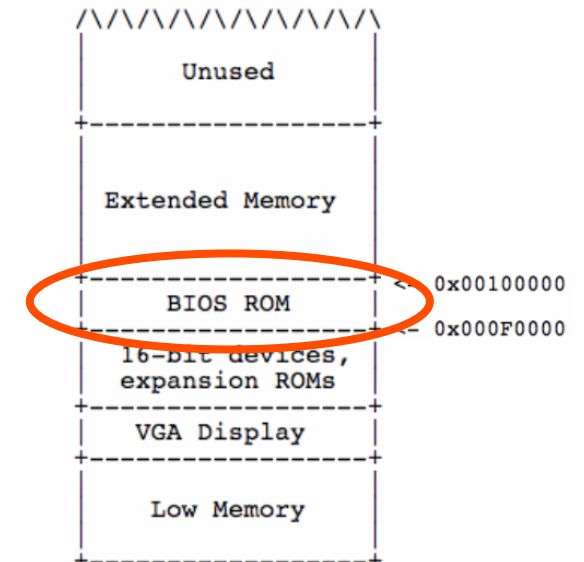
# Powering up : Reset



# Powering up : BIOS

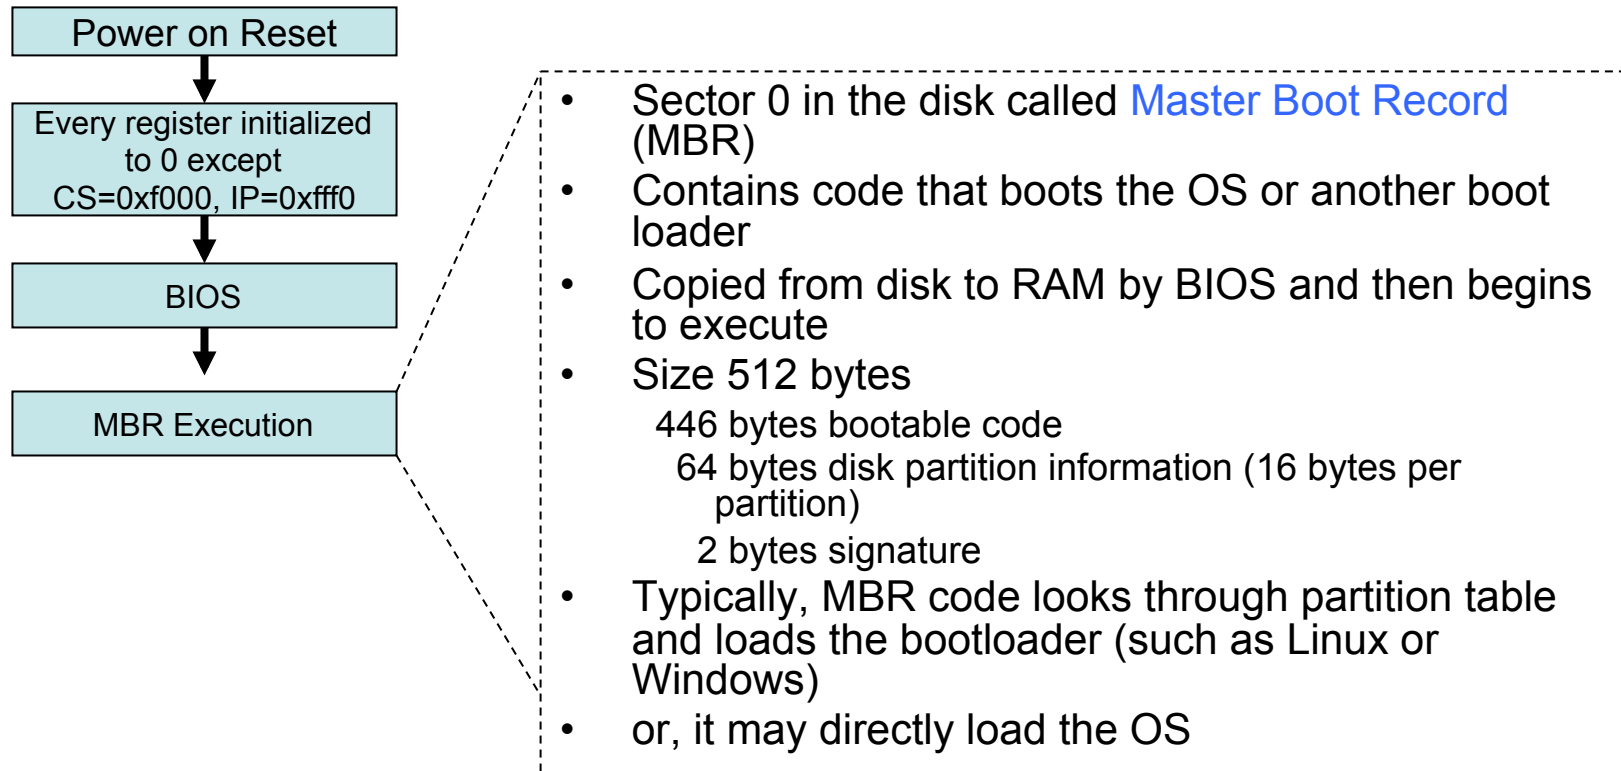


- Present in a small chip connected to the processor
  - Flash/EPROM/E<sup>2</sup>PROM
- Does the following
  - Power on self test
  - Initialize video card and other devices
  - Display BIOS screen
  - Perform brief memory test
  - Set DRAM memory parameters
  - Configure Plug & Play devices
  - Assign resources (DMA channels & IRQs)
  - Identify the boot device
    - Read sector 0 from boot device into memory location 0x7c00
    - Jumps to 0x7c00






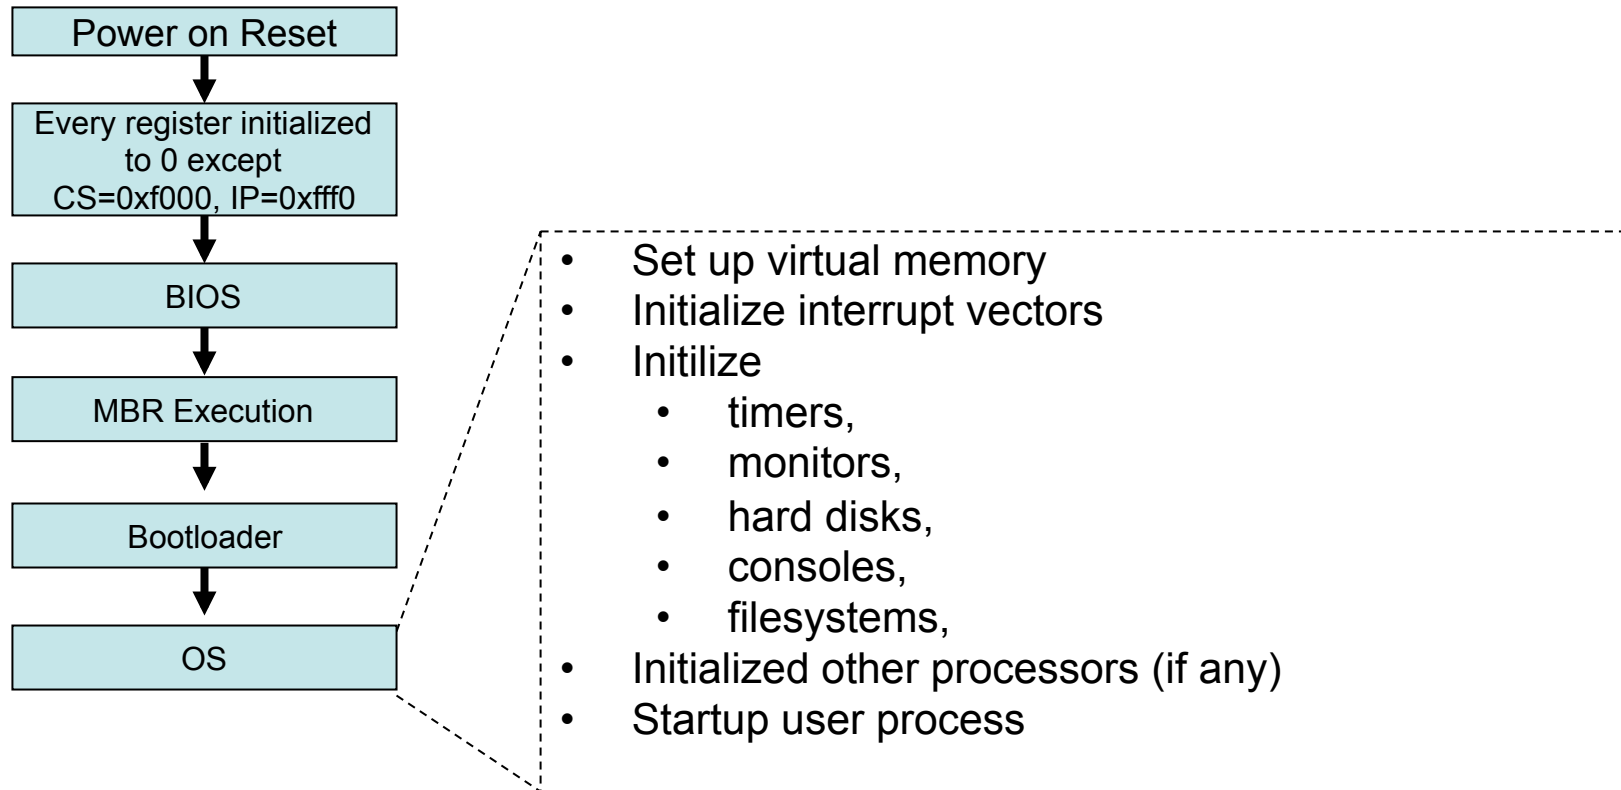
# Powering up : MBR





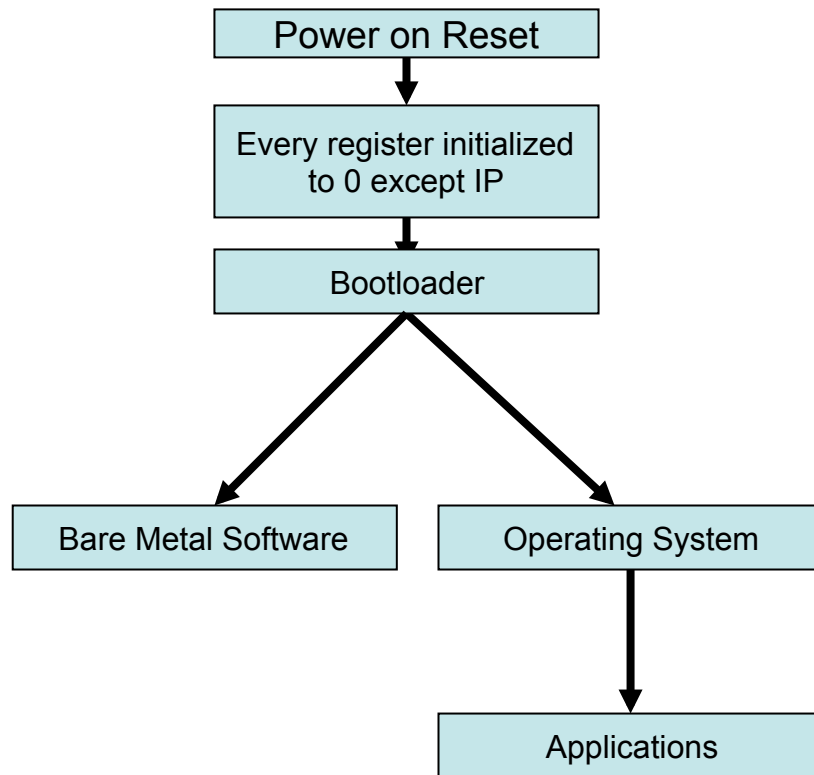
# Powering Up : bootloader



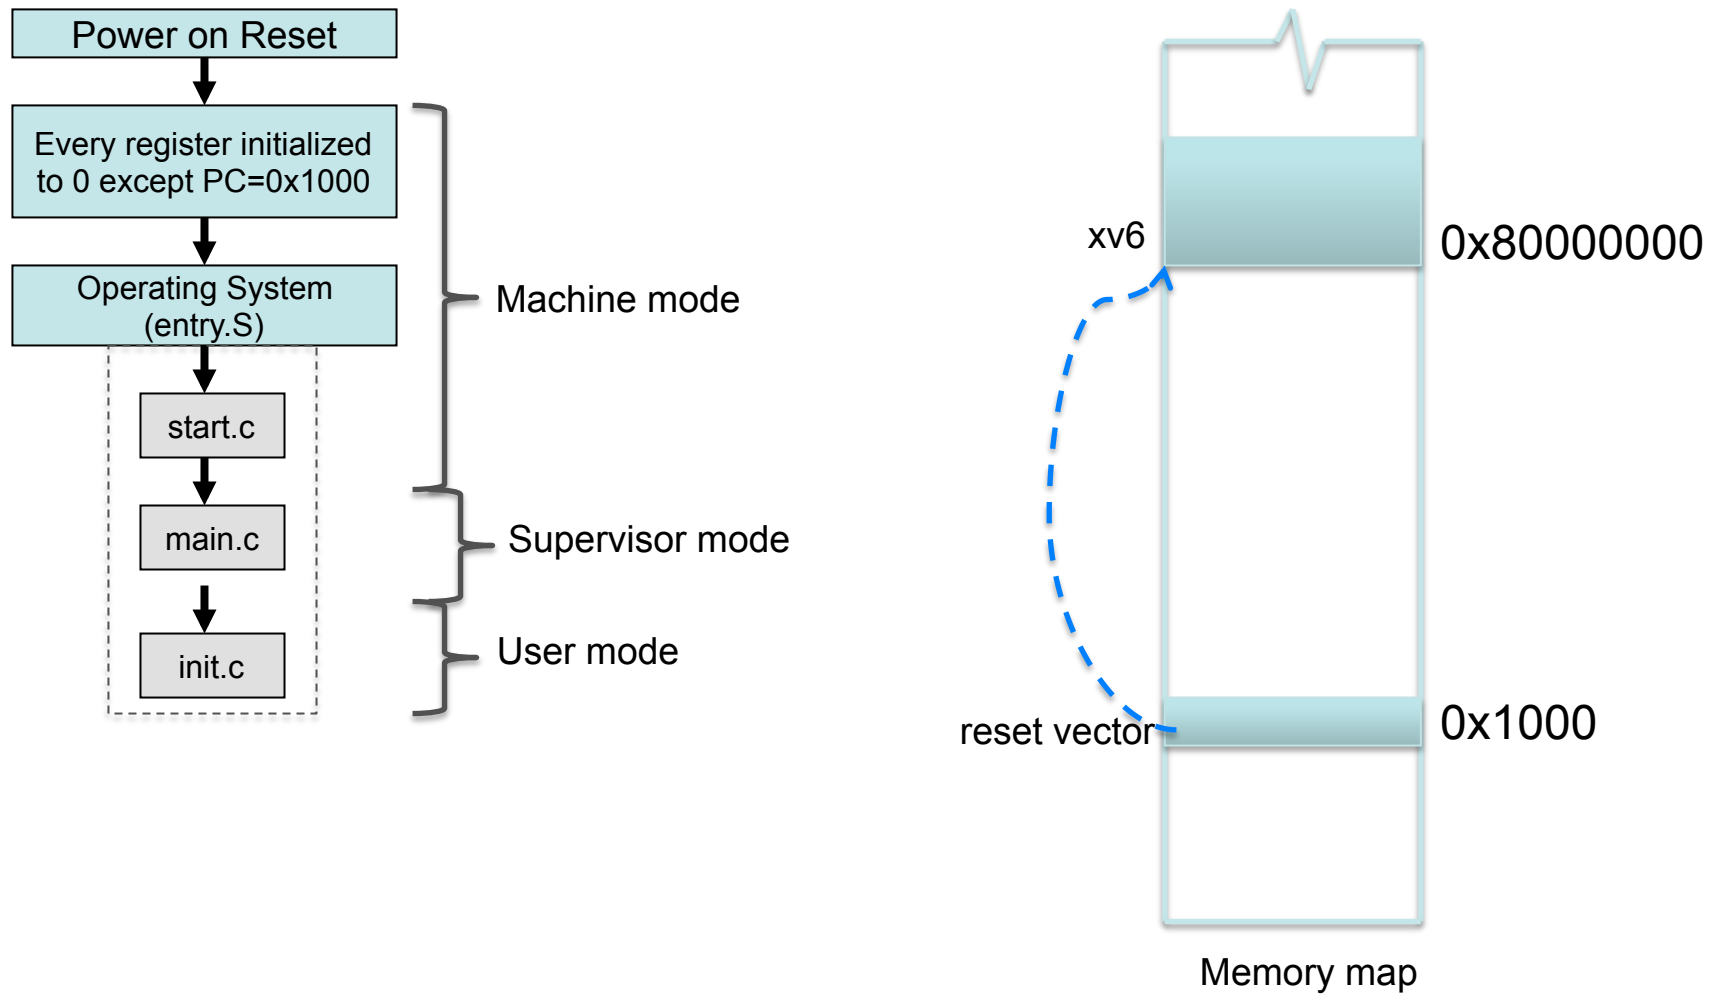
# Powering Up : OS



# Powering up an Embedded Device



# Powering up xv6 on qemu



# Linker Descriptor

## kernel.ld

```
OUTPUT_ARCH( "riscv" )
ENTRY( _entry )
```

### SECTIONS

```
{
    /*
     * ensure that entry.S / _entry is at 0x80000000,
     * where qemu's -kernel jumps.
     */
    . = 0x80000000;
    .text :
    {
        *(.text)
        . = ALIGN(0x1000);
        *(trampsec)
    }

    . = ALIGN(0x1000);
    PROVIDE(etext = .);

    /*
     * make sure end is after data and bss.
     */
    .data : {
        *(.data)
    }
    .bss : {
        *(.bss)
        *(.sbss*)
        PROVIDE(end = .);
    }
}
```

First function to be executed (before main)


Kernel code to be placed in 0x80000000

An indicator of the end of kernel code

After the kernel code, place data and then bss


An indicator of the end of kernel

# Powering Up : xv6 on qemu




- Memory Symmetry
  - All processors in the system share the same memory space
  - Advantage : Common operating system code
- I/O Symmetry
  - All processors share the same I/O subsystem
  - Every processor can receive interrupt from any I/O device

# Powering Up : xv6 on qemu



`csrr a1, mhartid` ; instruction that reads the hartid

All cores start the same way with IP=0x1000, and then jumping to 0x80000000






# xv6-boot

## entry.S

```
6 .section .data
7 .globl stack0
8 .section .text
9 .globl start
10 .section .text
11 .globl _entry
12 _entry:
13     # set up a stack for C.
14     # stack0 is declared in start.c,
15     # with a 4096-byte stack per CPU.
16     # sp = stack0 + (hartid * 4096)
17     la sp, stack0
18     li a0, 1024*4
19     csrr a1, mhartid
20     addi a1, a1, 1
21     mul a0, a0, a1
22     add sp, sp, a0
23     # jump to start() in start.c
24     call start
25 junk:
26     j junk
```



# xv6-boot

## entry.S

```
6 .section .data
7 .globl stack0
8 .section .text
9 .globl start
10 .section .text
11 .globl _entry
12 _entry:
13     # set up a stack for C.
14     # stack0 is declared in start.c,
15     # with a 4096-byte stack per CPU.
16     # sp = stack0 + (hartid * 4096)
17     la sp, stack0
18     li a0, 1024*4
19     csrr a1, mhartid
20     addi a1, a1, 1
21     mul a0, a0, a1
22     add sp, sp, a0
23     # jump to start() in start.c
24     call start
25 junk:
26     j junk
```

define label for stack0 (something like extern)  
actual stack0 defined in start.c

\_entry → will reside at 0x80000000

set sp to stack0 and a0 to 4KB

read the hartid for the processor core using  
the corresponding CSR

set sp appropriately for the core  
why increment a1?

call the C code (present in start.c)

Unreachable code

## start.c

# xv6-boot

```
19 // entry.S jumps here in machine mode on stack0.
20 void
21 start()
22 {
23     // set M Previous Privilege mode to Supervisor, for mret.
24     unsigned long x = r_mstatus();
25     x &= ~MSTATUS_MPP_MASK;
26     x |= MSTATUS_MPP_S;
27     w_mstatus(x);
28
29     // set M Exception Program Counter to main, for mret.
30     // requires gcc -mmodel=medany
31     w_mepc((uint64)main);
32
33     // disable paging for now.
34     w_satp(0);
35
36     // delegate all interrupts and exceptions to supervisor mode.
37     w_medeleg(0xffff);
38     w_mideleg(0xffff);
39
40     // ask for clock interrupts.
41     timerinit();
42
43     // keep each CPU's hartid in its tp register, for cpuid().
44     int id = r_mhartid();
45     w_tp(id);
46
47     // switch to supervisor mode and jump to main().
48     asm volatile("mret");
49 }
```

**Moves from  
machine mode  
to supervisor  
mode.**

## start.c

# xv6-boot

```
19 // entry.S jumps here in machine mode on stack0.
20 void
21 start()
22 {
23     // set M Previous Privilege mode to Supervisor, for mret.
24     unsigned long x = r_mstatus();
25     x &= ~MSTATUS_MPP_MASK;
26     x |= MSTATUS_MPP_S;
27     w_mstatus(x);
28
29     // set M Exception Program Counter to main, for mret.
30     // requires gcc -mmodel=medany
31     w_mepc((uint64)main);
32
33     // disable paging for now.
34     w_satp(0);
35
36     // delegate all interrupts and exceptions to supervisor mode.
37     w_medeleg(0xffff);
38     w_mideleg(0xffff);
39
40     // ask for clock interrupts.
41     timerinit();
42
43     // keep each CPU's hartid in its tp register, for cpuid().
44     int id = r_mhartid();
45     w_tp(id);
46
47     // switch to supervisor mode and jump to main()
48     asm volatile("mret");
49 }
```

previous mode = supervisor

Delegate all interrupts to be handled by Supervisor

Goto "previous mode" at the address "previous PC"



# xv6-boot

```
9 // start() jumps here in supervisor mode on all CPUs.
10 void
11 main()                main.c
12 {
13     if(cpuid() == 0){
14         consoleinit();
15         printfinit();
16         printf("\n");
17         printf("xv6 kernel is booting\n");
18         printf("\n");
19         kinit();        // physical page allocator
20         kvminit();       // create kernel page table
21         kvminithart();   // turn on paging
22         procinit();      // process table
23         trapinit();      // trap vectors
24         trapinithart();  // install kernel trap vector
25         plicinit();      // set up interrupt controller
26         plicinithart();  // ask PLIC for device interrupts
27         binit();         // buffer cache
28         iinit();         // inode cache
29         fileinit();      // file table
30         virtio_disk_init(); // emulated hard disk
31         userinit();      // first user process
32         __sync_synchronize();
33         started = 1;
34     } else {
35         while(started == 0)
36             ;
37         __sync_synchronize();
38         printf("hart %d starting\n", cpuid());
39         kvminithart();    // turn on paging
40         trapinithart();   // install kernel trap vector
41         plicinithart();   // ask PLIC for device interrupts
42     }
43
44     scheduler();
45 }
```

**Core specific initialization  
(done by all cores)**

**System specific initialization  
(done only by core 0)**

Executed by all cores  
except core0 after started=1

# xv6-boot

Core specific initialization  
(done by all cores)

System specific initialization  
(done only by core 0)

```
9 // start() jumps here in supervisor mode on all CPUs.
10 void
11 main()
12 {
13     if(cpuid() == 0){
14         consoleinit();
15         printfinit();
16         printf("\n");
17         printf("xv6 kernel is booting\n");
18         printf("\n");
19         kinit();           // physical page allocator
20         kvminit();         // create kernel page table
21         kvminithart();     // turn on paging
22         procinit();        // process table
23         trapinit();        // trap vectors
24         trapinithart();    // install kernel trap vector
25         plicinit();        // set up interrupt controller
26         plicinithart();    // ask PLIC for device interrupts
27         binit();           // buffer cache
28         iinit();           // inode cache
29         fileinit();        // file table
30         virtio_disk_init(); // emulated hard disk
31         userinit();        // first user process
32         __sync_synchronize();
33         started = 1;
34     } else {
35         while(started == 0)
36             ;
37         __sync_synchronize();
38         printf("hart %d starting\n", cpuid());
39         kvminithart();     // turn on paging
40         trapinithart();    // install kernel trap vector
41         plicinithart();    // ask PLIC for device interrupts
42     }
43
44     scheduler();
45 }
```